

Beacon Markets · 3-Minute Demo Runbook

Snowflake Summit · Pediment Booth · for the recorded video walkthrough of the Beacon Markets dbt Wizard demo. Looped unattended on the booth screen.

MANUAL MODEL BUILD
TODAY

5–8 days

DBT WIZARD BUILD

92 sec

VAR EXPOSURE
UNMODELED

\$184M

UNTIL RISK COMMITTEE

14 hr

ODI architecture · sources, lakehouse, consumers

Fivetran FIX gateway, market-data feed handlers, clearinghouse, CFTC EDGAR, FINRA, Salesforce member CRM, reference data	Iceberg + MDLS Open lakehouse on S3, Snowflake Horizon catalog, ACID snapshots, time travel	dbt Labs Bronze → Silver → Gold → Platinum marts, committed as Iceberg snapshots	dbt Wizard Build-time AI · 92s authoring of governed gold marts	Snowflake + Athena + Trino Same Iceberg bytes, no copies, external catalogs
--	---	--	---	---

One set of files on S3. Three engines read the same Parquet through the Snowflake Horizon catalog. Governed end-to-end: CFTC Reg AT · SRO surveillance rules · row-level access by desk.

DATA ENGINEERS ANCHOR HERE

ODI Architecture · dbt-wizard · Pipeline · Activations

Show the loop, the four sub-agents, the materialization, the connector deep-links. This is your DE proof.

BUSINESS USERS ANCHOR HERE

Trading Floor, Markets, Surveillance, Risk & Clearing, Regulatory

Show the answers, the KPI tiles, the ticker, the desk read. The same governed gold layer powers all of it.

Frame the demo on business outcomes — never functionality

PAIN	REASON	IMPACT	CAPABILITY	PROOF
Chief Risk Officer can't get intraday OI concentration in time for the Risk Committee	Every new cut = 5–8 day DE backlog	\$184M of house VaR exposure left unmodeled, Risk Committee meets blind	Build-time AI on the governed gold layer	Live model build in 92 seconds — on screen

The mantra: the win is not that dbt Wizard can author a model. The win is that the Chief Risk Officer walks into the next Risk Committee meeting with a tested, lineage-tracked answer that did not exist this morning. Every beat below ties to that one sentence.

The 3-minute script · five beats

00:00 – 00:25

Beat 1 · The stakeholder hook

WHAT YOU SAY (25S)	WHAT YOU SHOW
"It's Friday morning at Beacon Markets. The Chief Risk Officer needs intraday open-interest concentration by member, by product, in 15-minute buckets, across the rates complex — for the last five trading days. The Risk Committee meets in 14 hours. Today, an answer like this costs five to eight days of data-engineering work and a missed meeting. We're about to change that."	Open the Trading Floor (/). The CRO's question and the dbt-wizard callout are the first thing in frame. Outcome → Anchors the audience on a real business question, not a feature.

00:25 – 00:50

Beat 2 · The cost of the gap

WHAT YOU SAY (25S)	WHAT YOU SHOW
"There's no gold-layer mart that answers this. Five to eight days from now, after a request, a backlog, a scope, a build, and a peer review, the engineering team might land it. The Risk Committee will have already met. \$184 million of house VaR exposure stays unmodeled in the meantime — and that's just the exposure we can see."	Stay on the Trading Floor. Highlight the dbt-wizard callout tile: 92s vs 5–8 days, \$184M at risk. Outcome → Quantifies the cost of the gap. Sets up the 92-second proof.

00:50 – 02:20

Beat 3 · The live build (proof moment · DE-anchor)

WHAT YOU SAY (90S)	WHAT YOU SHOW
"We click See the live build. Four sub-agents — Explorer, Summary, Worker, Verification — open the dbt project, find the right grain, pull a slice of the silver layer, write the new model, write its tests, and materialize it as an Iceberg snapshot. The viewer sees the SQL appear character by character. By the time we're done talking, the model exists, the tests pass, the lineage graph has updated."	Click into /dbt-wizard . Let the playback run at 1x. The step rail moves Discovery → Schema understanding → Data inspection → Model creation → Test authoring → Materialization. Terminal-window aesthetic. Outcome → Data engineers see the loop end-to-end. The wow is that the output is a tested, governed asset — not a snippet pasted into Slack.

Pause once at: "Worker is deriving the 15-minute bucket from event_ts — not a spreadsheet-grade approximation." And again at: "Verification is writing a uniqueness test on member × product × bucket. This is a governed gold asset, not a SQL snippet."

02:20 – 02:45

Beat 4 • The outcome (P&L tie-back • business-anchor)

WHAT YOU SAY (25S)	WHAT YOU SHOW
"The new mart is materialized. The Risk Committee convenes in 14 hours, not eight days. The Chief Risk Officer walks in with the answer: three clearing members are carrying 61% of rates open interest, concentrated in the last two hours of trading — a margin-call risk, not a six-month capacity build. And the next time anyone asks this question — energy, agriculture, equities, anywhere — the model is already there."	Cut to the /dbt-wizard outcome panel — 92s, tests passing, downstream consumers listed. Then quick scroll through the business-user pages (Markets, Surveillance, Risk & Clearing, Regulatory) so the audience sees the same gold layer powering the desk view. Outcome → Closes the loop from question to action AND shows the business-user surface running on the same data.

02:45 – 03:00

Beat 5 • The architectural ask

WHAT YOU SAY (15S)	WHAT YOU SHOW
"Fivetran lands every source — FIX, market data, clearing, CFTC, FINRA, member CRM. Iceberg on S3 holds the bytes; the Snowflake Horizon catalog governs the snapshots. dbt Labs builds the marts. dbt Wizard authors the gaps. Snowflake, Athena, and Trino all read the gold layer — and so does NewCo Activations, pushing critical surveillance alerts straight into case management. One set of files, three engines, zero manual re-keying. Want to see how this maps to your gold layer? Stop by the Pediment booth or scan the QR."	Final card with the architecture loop, QR code to fivetran.com/contact. Tag dbt Labs co-brand. Outcome → Calls out the architectural pattern. Invites a discovery conversation.

Business outcomes to weave through — never lead with features

OUTCOME THE BUYER FEELS	WHAT IT REPLACES	HOW TO PHRASE IT
Risk Committee gets the answer same-day	5–8 day DE backlog	In time for the meeting, not a week late
\$184M VaR exposure sized in 92 seconds	Open gap window	The cost of the gap closed before lunch
Reusable governed mart for every future ask	One-off notebook scripts	The next person who asks gets it for free
CDO does not add headcount to keep up	Hiring DE backfill	Your team stops being the bottleneck
Auditor sees lineage + tests, not a query history	Manual evidence packs	Examiner-ready by default

Anti-patterns · what NOT to say

Don't say: "It runs dbt show against the warehouse and authors a new model file." **Say:** "Ninety-two seconds from question to materialized answer."

Don't say: "Cortex Analyst." Cortex is intentionally absent from this demo. The story is build-time AI on the governed lake — the in-app feature is called dbt Wizard.

Don't say: "dbt Wizard generates SQL." **Say:** "dbt Wizard authors a governed asset — schema contract, tests, lineage, ownership."

Don't say: "Look at the typing animation." The animation is texture, not the value.

Don't open with the architecture diagram. Open with the Chief Risk Officer's question. Architecture is the closer.

Don't skip the business pages. After the wizard finishes, scroll through Markets, Surveillance, Risk & Clearing, Regulatory so the audience sees the persona surface that runs on the same gold layer.

Discovery prompts to anchor before the demo

BEFORE YOU PRESS PLAY

- "What questions does your Chief Risk Officer ask that your team cannot answer intraday?"
- "When a desk asks for a new cut of position or surveillance data, what's the average time from ask to merged dbt model?"
- "How many one-off SQL scripts are sitting in someone's Snowflake history that never made it back into the governed layer?"
- "Who is your Economic Buyer for new data investment? Likely audience: CRO, CDO, Chief Compliance Officer, or Head of Surveillance."

AFTER THE DEMO LANDS

- "Which of those agent steps — Discovery, Schema understanding, Data inspection, Model creation, Test authoring, Materialization — is where your team spends the most time today?"
- "If we sized a pilot on three of your highest-value recurring Risk Committee questions, what would success look like in 90 days?"
- "Who needs to be in a 30-minute architecture review to see how this lands on your gold layer? Audience: CRO, CDO, Chief Compliance Officer, or Head of Surveillance."

Personas at the booth · split the time

DATA ENGINEER · 60% OF BOOTH TIME

Trading Floor → dbt-wizard → ODI Architecture → Pipeline

Lead with the question, show the four sub-agents, then walk the architecture (Fivetran + Iceberg/MDLS + dbt Labs + dbt Wizard + Snowflake/Athena/Trino) and the Pipeline page with real Fivetran connector deep-links.

BUSINESS USER · 40% OF BOOTH TIME

Markets, Surveillance, Risk & Clearing, Regulatory

After the wizard finishes, drop into the desk pages. Show that the same governed gold layer powers the trading-floor read. The build-time AI is invisible to them — they see the answer.

Required follow-up materials

- One-pager · Beacon Markets · at fivetran-jasonchletsos.github.io/Fivetran-Demo-Repository/one-pager/capital-markets/
- Live demo URL · fivetran-jasonchletsos.github.io/CapitalMarkets-ODI-Demo/
- 3-minute walkthrough video (this recording)
- Architecture diagram on the ODI Architecture tab of the demo
- Cross-industry credibility · Pediment Bank

Booth-floor reminder: the demo is built to run unattended on the booth screen. If a prospect walks up mid-playback, do not restart. Pick up the narration where the playback is — that is also the point. The build finishes whether someone is watching or not. That is the entire pitch.

The gold layer gets acted on, not just queried

Activations is the fourth native Fivetran stage, right after Transformations. The instant `gold.fct_alert_evidence` writes a row where `severity = 'critical'`, `pattern_confidence_score ≥ 0.85`, and the member is already on the members-under-review list — a corroborated, repeat high-confidence pattern, not a one-off flag — Activations pushes the alert straight into NICE Actimize Case Manager and auto-opens a fully populated investigation case. No reverse-ETL vendor, no second copy of the data, no analyst re-keying.

TRIGGER

confidence ≥ 0.85

critical + repeat member under review

DESTINATION

Actimize

Case Manager · auto-opened case

TIME TO OPEN CASE

12 min → <30 sec

18 critical alerts / day

WHAT YOU SAY	WHAT YOU SHOW
"Halcyon FCM just tripped layered-cancels spoofing in ZB futures — third critical alert this week, confidence score 0.91. Today an analyst opens Actimize, keys in the alert ID, member ID, order IDs, pastes an evidence link — call it 12 minutes. Times 18 critical alerts a day, that's 3-plus analyst-hours gone before anyone investigates anything. Watch what happens the second <code>gold.fct_alert_evidence</code> writes this row."	Live tail on <code>gold.fct_alert_evidence</code> — Halcyon FCM row lands: severity critical, confidence 0.91, member already under review.
"Cut to the Activations sync log — same row, timestamped four seconds later."	Architecture page → Activations panel → /activations-live: Segment Definition → Field Mapping → Sync Preview → API Push.
"Cut to Actimize Case Manager — case ALT-2026-58841 is already open: subject Halcyon FCM, risk score 91, evidence attached, typology 'spoofing — layered cancels.'"	/activations-live → Destination Confirmation table, case card rendered with the matching member, risk score, and typology fields.
"Zero manual entry, zero re-typed order IDs, and the audit trail proves the case opened within seconds of detection — not whenever an analyst got to it."	Before/After framing: 12-minute manual re-keying, once per critical alert, replaced by one governed sync, zero code.

FIELD MAPPING · GOLD.FCT_ALERT_EVIDENCE → ACTIMIZE CASE MANAGER

SOURCE COLUMN	DESTINATION FIELD
alert_id	case_external_id
member_id	subject_party_id
pattern_confidence_score	risk_score
evidence_uri	evidence_attachment_url
pattern_type	typology_code

Current snapshot: 18 critical alerts / day. Closing the critical-alert-to-open-case gap reclaims roughly 3.6 analyst-hours daily for investigation instead of data entry, and closes the documented-escalation gap exchanges must show under CFTC Reg AT and SRO surveillance rules. Synthetic data — not a real exchange.